

University of Central Florida
Department of Computer Science
CDA 5106: Fall 2021
Machine Problem 3: Hybrid Predictor

1. Ground rules

- i. All students must work alone.
- ii. Sharing of code between students is considered cheating and will receive appropriate action in accordance with University policy. The TAs will scan source code through various tools available to us for detecting cheating. Source code that is flagged by these tools will be dealt severely: 0 on the project and referral to the Office of Student Conduct for sanctions.
- iii. You must do all your work in C/C++ or Java language. Exceptions must be approved. The C language is fine to use as that is what many students are trained in. Basic C++ extensions to the C language (classes instead of structs) are encouraged (but by no means required) because it enables more straightforward code reuse.
- iv. Use of Linux environment is required. This is the platform where the TAs will compile and test your simulator.

2. Project Description

In this project, you will construct a simulator for a hybrid predictor using two gshare predictors: bimodal($n=0$) and gshare($n>0$) and use it to design a branch predictor well suited to the SPECint95 benchmarks.

3. Simulator Specification

3.1 Hybrid Branch Predictor

Model a **hybrid predictor** that selects between a *bimodal* and a *gshare* predictor, using a chooser table of 2^k 2-bit counters, where k is the number of low order PC bits used to index the chooser table. All counters in the chooser table are initialized to **1** at the beginning of the simulation. You can reuse your source code from Machine Problem 2 for bimodal and gshare predictors. For reference, instructions for bimodal and gshare predictors are included in section 3.2.

Steps:

When you get a branch from the trace file, there are six top-level steps:

- (1) Obtain two predictions, one from *gshare* predictor (follow steps 1 and 2 in section 3.2.2) and one from *bimodal* predictor (follow steps 1 and 2 in section 3.2.1).
- (2) Determine the branch's **index** into the chooser table. The index for the chooser table is bit $k+1$ to bit **2** of the branch PC (i.e., as before, discard the lowest two bits of the PC).
- (3) Make an overall prediction. Use **index** to get the branch's chooser counter from the chooser table. If the chooser counter value is greater than or equal to **2**, then use the prediction that was obtained from the *gshare* predictor, otherwise use the prediction that was obtained from the *bimodal* predictor.

- (4) Update the selected branch predictor based on branch's actual outcome. Only the branch predictor that was selected in step 3, above, is updated (if *gshare* was selected, follow step 3 in Section 3.1.2, otherwise follow step 3 in section 3.1.1).
- (5) Note that the *gshare*'s global branch history register must always be updated, even if *bimodal* was selected (follow step 4 in section 3.1.2).
- (6) Update the branch's chooser counter using the following rule:

	Results from predictors:		
	<i>both incorrect or both correct</i>	<i>gshare correct, bimodal incorrect</i>	<i>bimodal correct, gshare incorrect</i>
Chooser counter update policy:	no change	increment (but saturates at 3)	decrement (but saturates at 0)

3.2 GShare Branch Predictors

The *gshare* branch predictor takes parameters $\{m, n\}$, where:

- m is the number of low-order PC bits used to form the prediction table index. **Note:** discard the lowest two bits of the PC, since they are always zero, i.e., use bits $m+1$ through 2 of the PC.
- n is the number of bits in the global branch history register. **Note:** $n \leq m$. **Note:** n may equal zero, in which case we have the simplest *bimodal* branch predictor.

3.2.1. $n=0$: bimodal branch predictor

When $n=0$, the *gshare* predictor reduces to a simple *bimodal* predictor. In this case, the index is based on only the branch's PC, as shown in **Figure 1**.

Entry in the prediction table

An entry in the prediction table contains a single 3-bit counter. All entries in the prediction table should be initialized to 4 ("weakly taken") when the simulation begins.

Regarding branch interference

Different branches may index the same entry in the prediction table. This is called "interference". Interference is not explicitly detected or avoided: it just happens. (There is no tag array, no tag checking, and no "miss" signal for the prediction table!)

Steps:

When you get a branch from the trace file, there are three steps:

- (1) Determine the branch's **index** into the prediction table.
- (2) Make a prediction. Use **index** to get branch's counter from the prediction table. If the counter value is greater than or equal to 4, then the branch is predicted *taken*, else it is predicted *not-taken*.
- (3) Update the branch predictor based on the branch's actual outcome. The branch's counter in the prediction table is incremented if the branch was taken, decremented if the branch was not taken. The counter *saturates* at the extremes (0 and 7), however.




Figure 1. Bimodal branch predictor.

3.2.2. $n > 0$: gshare branch predictor

When $n > 0$, there is an n -bit global branch history register. In this case, the index is based on both the branch's PC and the global branch history register, as shown in **Figure 2**. The global branch history register is initialized to all zeroes (00...0) at the beginning of the simulation.

Steps:

When you get a branch from the trace file, there are three steps:

- (1) Determine the branch's **index** into the prediction table. Figure 2 shows how to generate the index: the current n -bit global branch history register is XORed with the lowermost n bits of the index (PC) bits.

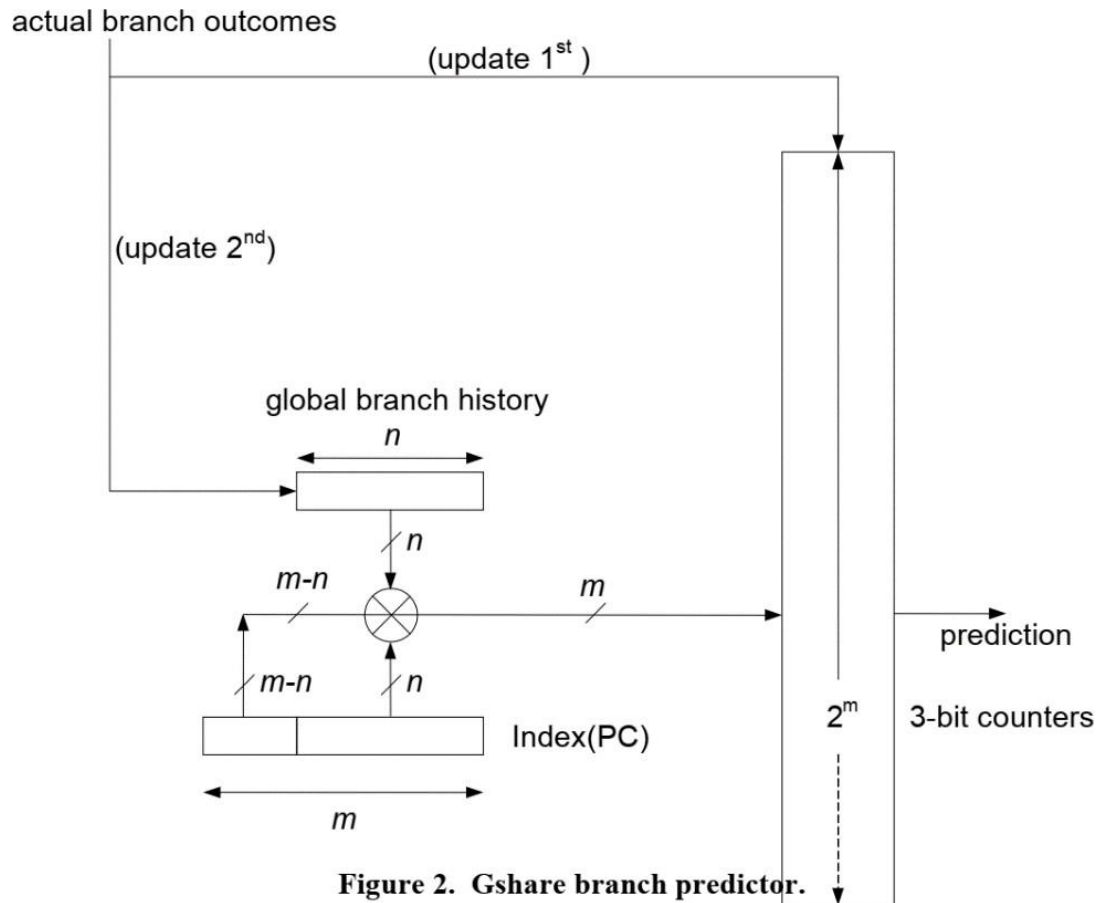
m -bit index = concatenate ($m-n$ MSB bits of lowermost m bits from PC,
XOR(n -bit global branch history, LSB n bits of PC))

Example: Let PC = 302d28, $m=9$, $n=3$, Global History register value: 000

Steps to get branch's index into the prediction table

1. Remove last 2 bits of PC as mentioned in section 3.2 ==> 0011 0000 0010 1101 0010 1000
 2. Get the lowermost m bits from the above value ==> 101 0010 10
 3. $m-n$ MSB bits of the above value ==> 101 001
 4. n bits of Global History register ==> 000
 5. n LSB bits of PC ==> 010
 6. XOR values from 4 and 5 ==> 010
 7. Concatenate values from 3 and 6 to get the m -bit index ==> 101 001 010
- (2) Make a prediction. Use **index** to get the branch's counter from the prediction table. If the counter value is greater than or equal to 4, then the branch is predicted *taken*, else it is predicted *not-taken*.
 - (3) Update the branch predictor based on the branch's actual outcome. The branch's counter in the prediction table is incremented if the branch was taken, decremented if the branch was not taken. The counter *saturates* at the extremes (0 and 7), however.

- (4) Update the global branch history register. Shift the register right by 1 bit position and place the branch's actual outcome into the most significant bit position of the register.



4. Inputs to Simulator

The simulator reads a trace file in the following format:

```
<hex branch PC> t|n  
<hex branch PC> t|n  
...
```

Where

- <hex branch PC> is the address of the branch instruction in memory. This field is used to index the predictors.
- "t" indicates the branch is actually taken (Note! *Not* that it is predicted taken!). Similarly, "n" indicates the branch is actually not-taken.

Example:

```
00a3b5fc t  
00a3b604 t  
00a3b60c n  
...
```

5. Outputs from Simulator

The simulator outputs the following measurements after completion of the run:

- a. Number of accesses to the predictor (i.e., number of branches)
- b. Number of branch mispredictions (predicted *taken* when *not-taken*, or predicted *not-taken* when *taken*)
- c. Branch predictions rate (# of mispredictions / # of branches)

6. Validation and Other Requirements

6.1. Validation requirements

Sample simulation outputs will be provided on the webcourses. These are called "validation runs". You must run your simulator and debug it until it matches the validation runs.

Each validation run includes:

1. The branch predictor configuration
2. The final contents of the branch predictor
3. All measurements described in Section 5.

Your simulator must print outputs to the console (i.e., to the screen). (Also see Section 6.2 about the requirement.)

Your output must match both numerically and in formatting, because the TAs will literally “diff” your output with the correct output. You must confirm correctness of your simulator by following these two steps for each validation run:

1. Redirect the console output of your simulator to a temporary file. This can be achieved by placing `> your_output_file` after the simulator command.
2. Test whether or not your outputs match properly, by running this unix command:
`diff -i -w <your_output_file> <posted_output_file>`
The “-i -w” tags tell “diff” to treat upper-case and lower-case as equivalent and to ignore the amount of whitespaces between words. Therefore, you do not need to worry about the exact number of spaces or tabs as long as there is some whitespace where the validation runs have whitespace.

6.2. Compiling and running simulator

You will hand in source code and the TAs will compile and run your simulator. As such, you must meet the following strict requirements. Failure to meet these requirements will result in point deductions (see section “Grading”).

1. You must be able to compile and run your simulator on Linux machine. This is required so that the TAs can compile and run your simulator.
2. Along with your source code, you must provide a **Makefile** that automatically compiles the simulator. This Makefile must create a simulator named “sim”. The TAs should be able to type only “make” and the simulator will successfully compile. The TAs should be able to type only “make clean” to automatically remove object files and the simulator executable. An example Makefile is posted in the **MachineProblem3Fall2021.zip** folder, which you can copy and modify for your needs.
3. Your simulator must accept command-line arguments as follows:
To simulate a hybrid predictor: `sim hybrid <K> <M1> <N> <M2> <tracefile>`

where K is the number of PC bits used to index the chooser table, M1 and N are the number of PC bits and global branch history register bits used to index the gshare table (respectively), and M2 is the number of PC bits used to index the bimodal table.

(`<tracefile>` is the filename of the input trace.)

4. Your simulator must print outputs to the console (i.e., to the screen). This way, when a TA runs your simulator, he/she can simply redirect the output of your simulator to a filename of his/her choosing for validating the results.

6.3. Run time of simulator

Correctness of your simulator is of paramount importance. That said, making your simulator efficient is also important for a couple of reasons.

First, the TAs need to test every student's simulator. Therefore, we are placing the constraint that your simulator must finish a single run in **2 minutes** or less. If your simulator takes longer than 2 minutes to finish a single run, please see the TAs as they may be able to help you speed up your simulator.

Second, you will be running many experiments: many branch predictor configurations and multiple traces. Therefore, you will benefit from implementing a simulator that is reasonably fast.

[The following applies to students working with C.]

One simple thing you can do to make your simulator run faster is to compile it with a high optimization level. The example Makefile posted on the **MachineProblem3Fall2021.zip** folder includes the -O3 optimization flag.

Note that, when you are debugging your simulator in a debugger (such as gdb), it is recommended that you compile without -O3 and with -g. Optimization includes register allocation. Often, register-allocated variables are not displayed properly in debuggers, which is why you want to disable optimization when using a debugger. The -g flag tells the compiler to include symbols (variable names, etc.) in the compiled binary. The debugger needs this information to recognize variable names, function names, line numbers in the source code, etc. When you are done debugging, recompile with -O3 and without -g, to get the most efficient simulator again.

7. Tasks and Grading Breakdown

Match the three validation runs “val_hybrid_1.txt”, “val_hybrid_2.txt” and “val_hybrid_3.txt” posted on the webcourses for the HYBRID PREDICTOR. The TAs will also check that your simulator matches a mystery validation run.

40 points: Substantial Programming effort

Item	Points
Substantial simulator turned in	40

60 points: A working simulator with matched validation runs

Item	Points
val_hybrid_1.txt	15
val_hybrid_2.txt	15
val_hybrid_3.txt	15
Mystery run A	15

8. What to Submit via Webcourses

You must hand in a single zip file called **project3.zip**. Below is an example showing how to create **project3.zip** from an EOS Linux machine. Suppose you have a bunch of source code files (*.cc, *.h), and the Makefile.

zip **project3** *.cc *.h Makefile

project3.zip must contain the following (any deviation from the following requirements may delay grading your project and may result in point deductions, late penalties, etc.

1. **Source code.** You must include the commented source code for the simulator program itself. You may use any number of .cc/.h files, .c/.h files etc.
2. **Makefile:** See Section 6.2, item #2, for strict requirements. If you fail to meet these requirements, it may delay grading your project and may result in point deductions.

9. Penalties

Various deductions (out of 100 points):

Up to -10 points: for not complying with specific procedures. Follow all procedures very carefully to avoid penalties.

Cheating: Source code that is flagged by tools available to us will be dealt with according to University Policy. This includes a '0' for the project and other disciplinary actions.

